GRACE HOPPER CELEBRATION

ANITAB.ORG

Proof-of-concept to Production

How to scale your Deep Learning Models

Presenters: Swetha Mandava, Alex Qi

Contributor: Michael Mendelson

AGENDA

Deep Learning at Scale

- Why scaling
- Tips and Tricks

Desired outcomes

- Understand when to and how to scale
- Know the typical techniques to apply
- Understand the theory/concepts of typical scaling techniques

Deep Learning At Scale

Proof-of-concept

- Model architecture search
- Hyperparameter Optimization
- Toy dataset/Toy models
- Low frequency of retraining
- Non-optimized resource utilization

Production

- Models need to be trained and retrained with shorter times and higher frequency.
- Massive datasets
- Big & complex models

Pain Points

DL models training:

- **Time consuming**: can take days, weeks...
- **Capability:** is limited by memory capacity on batch size and model size

Scaling is necessary but hard:

- Convergence and Stability
- Computation and Scaling efficiency
- Hardware Limits

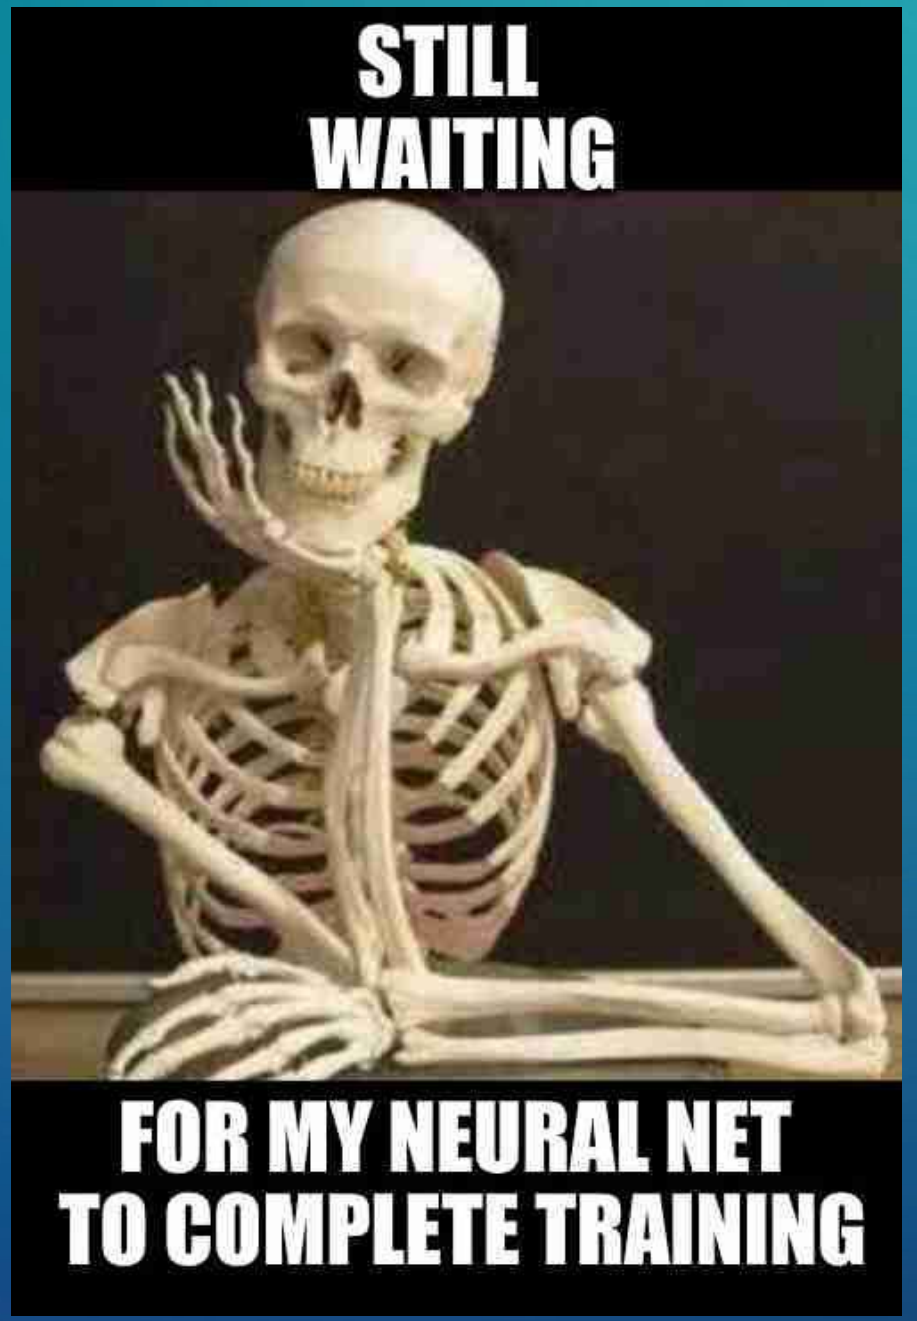

Src: https://mc.ai/ibm-offering-gpu-powered-notebooks/

Techniques for Scaling Deep Learning Training

- **Convergence and Stability**
  - Warmup
  - Linear Scaling Rule
  - LARS
- **Computation and Scaling efficiency**
  - Automatic Mixed Precision
- **Hardware limits on Dataset and Model Size**
  - User profiles/preference of millions of users
  - Data Parallelism
  - Model Parallelism

Scaling Success

BERT PreTraining on DGX SuperPOD	
# V100 GPUs	Time to train (hours)
16	58.4 (2.4 days)
256	3.9
1024	1.2
1472	0.9 (53 min)

GOAL :

1. Maintain Accuracy

2. Decrease Time to Train

Let's do it on a smaller model: NCF
BERT has 110 million parameters!

Recommender System: Neural Collaborative Filtering

Recommendation engines are everywhere…

- Personalized
- Convenient
- More efficient

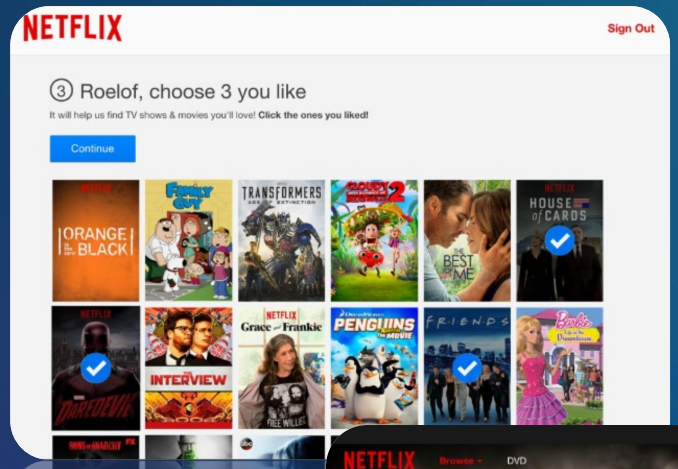

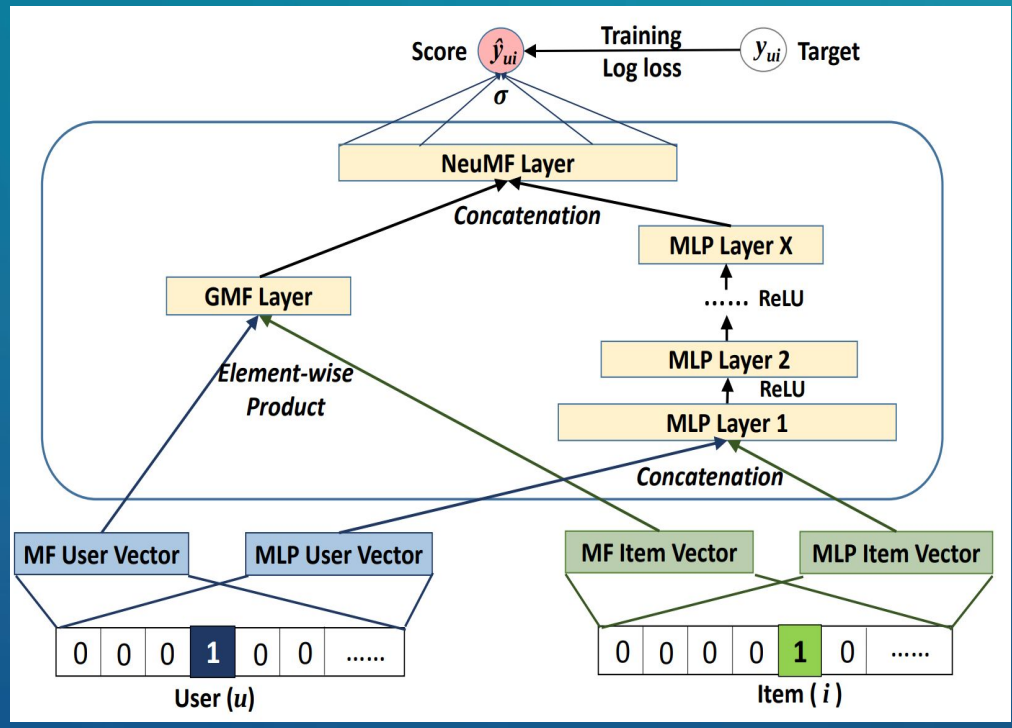

Deep Learning based Recommendation system model architecture

#GHC19

Launch Hands-on Task

- Create Account at https://courses.nvidia.com/join
- Go to courses.nvidia.com/dli-event
- Browser Recommendation: Chrome
- Use event code and create an account

Select the “Course” Tab

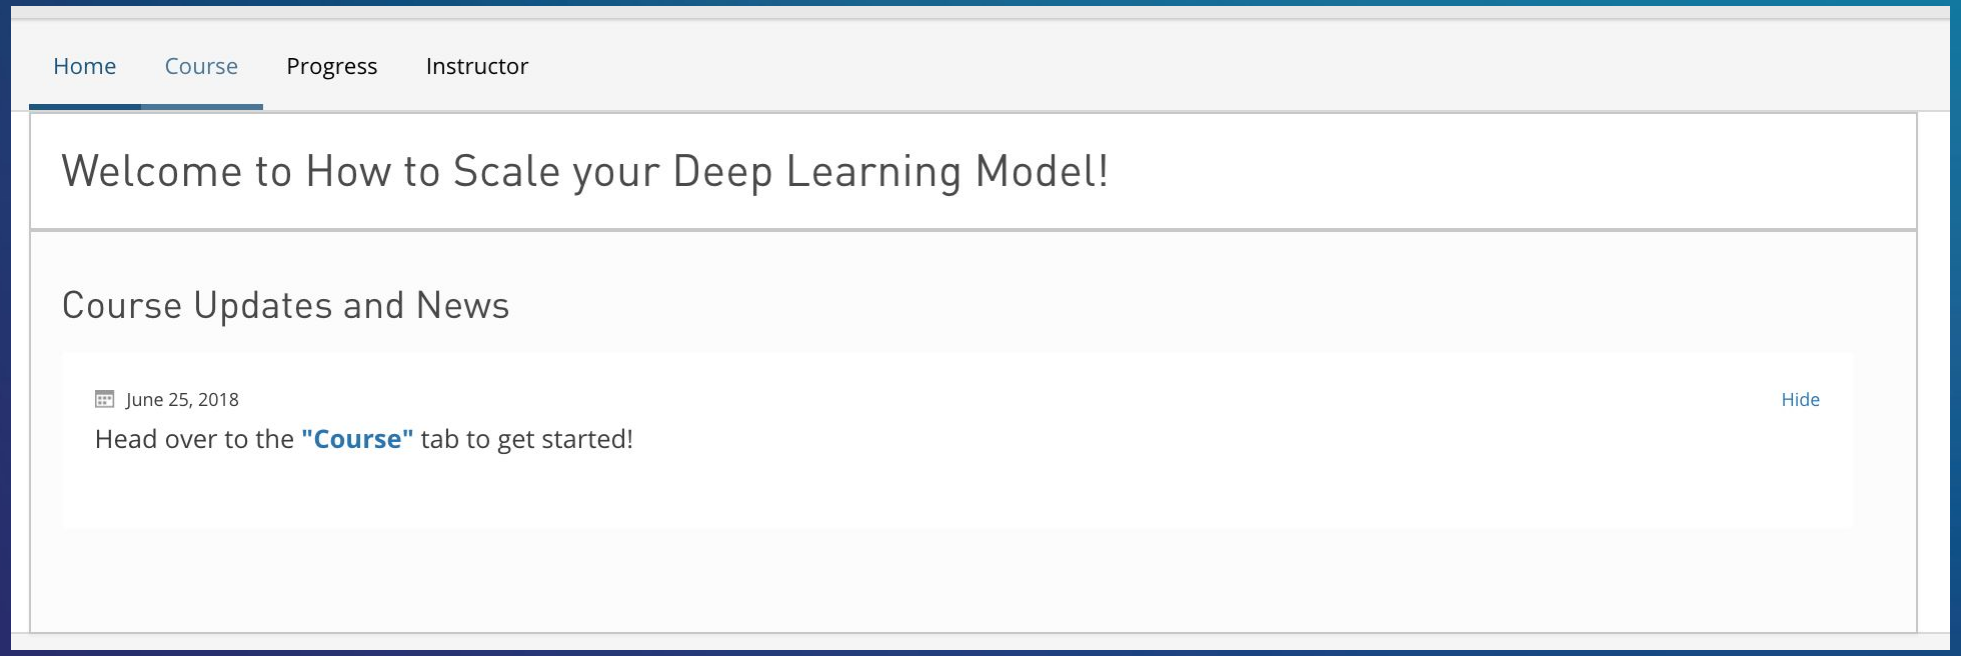

Open the first hands-on section

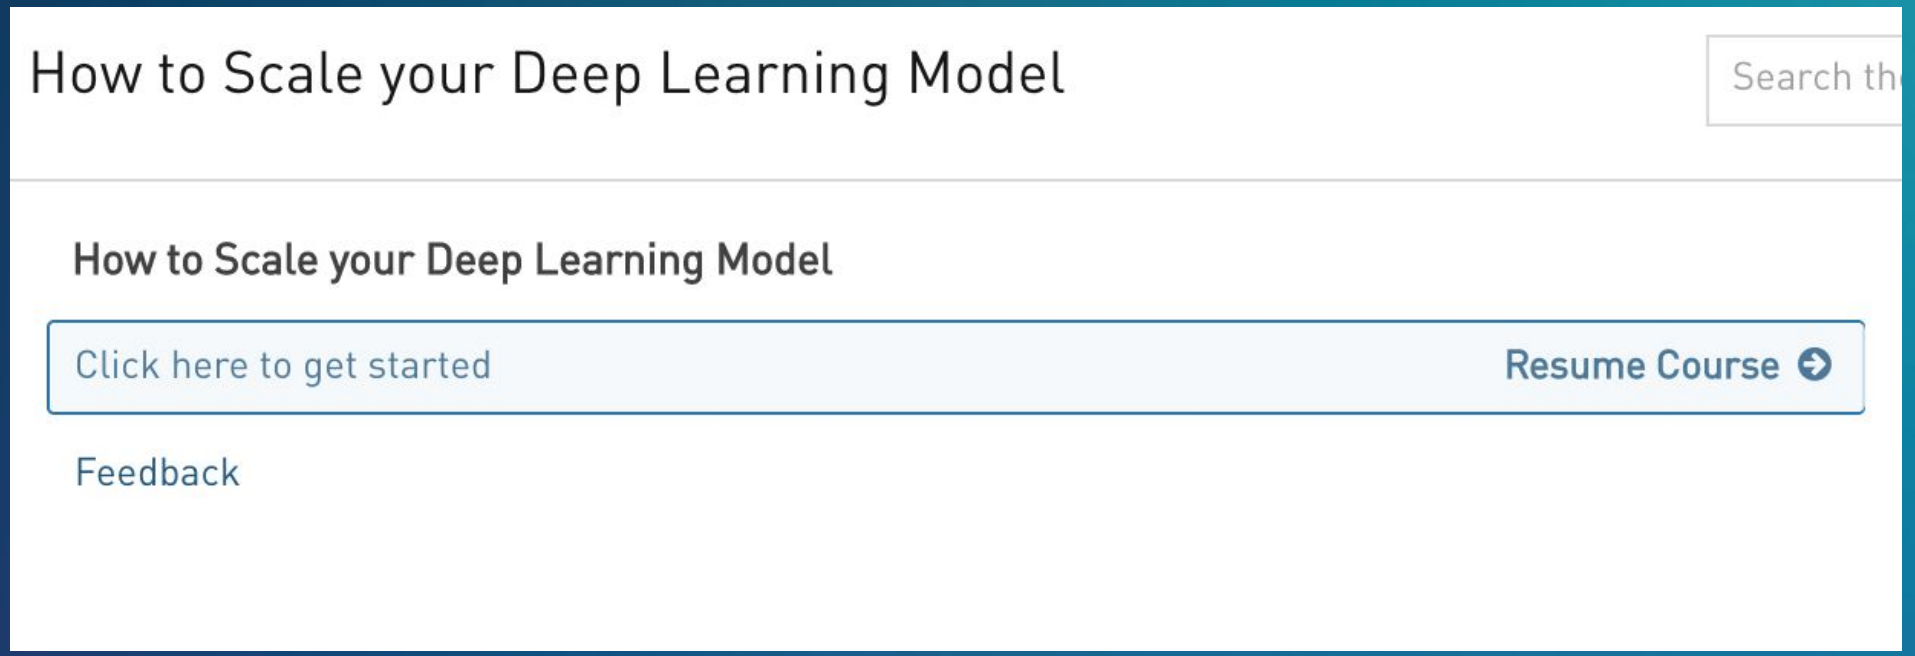

Select the Start button and wait

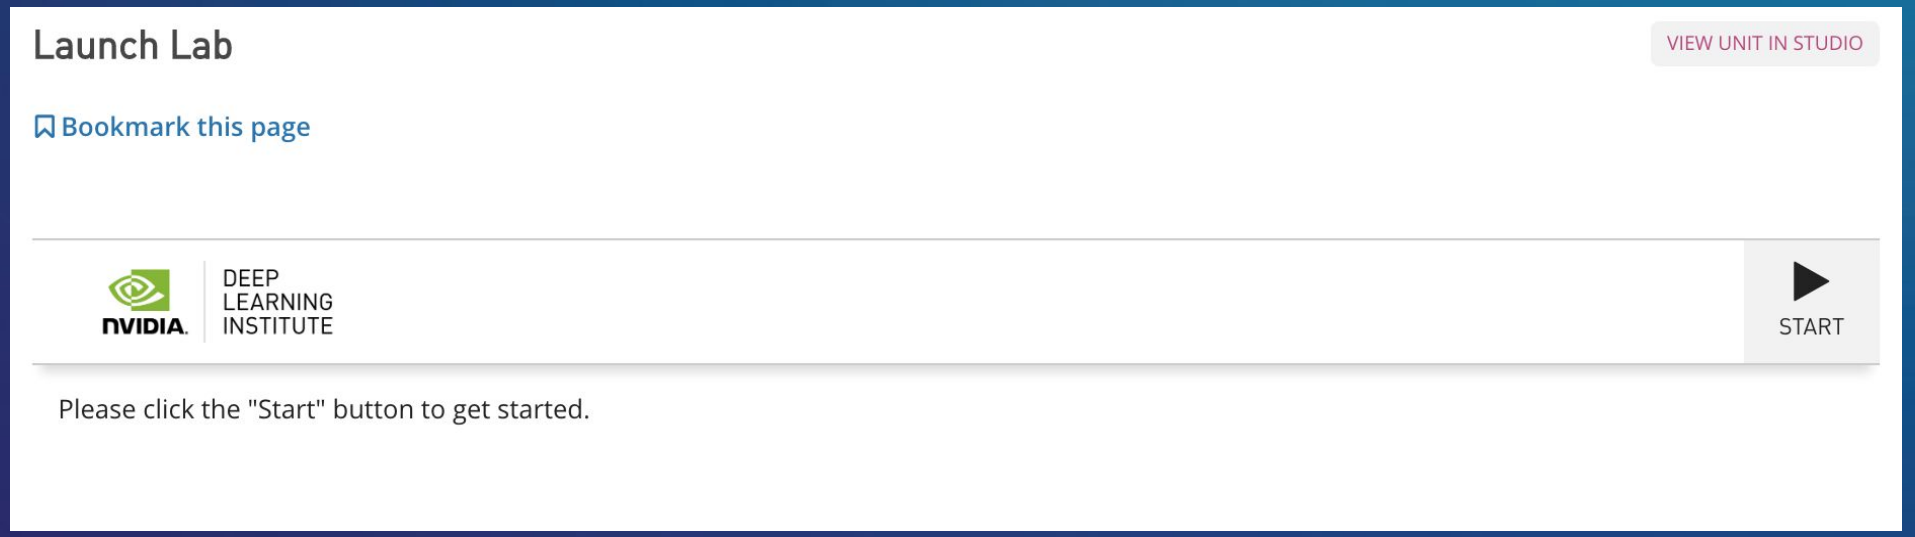

What else can we do?

- Is 192 still the maximum batch size we can use?
- Can we scale further than 1 GPU?
    - Data Parallelism

https://github.com/NVIDIA/DeepLearningExamples

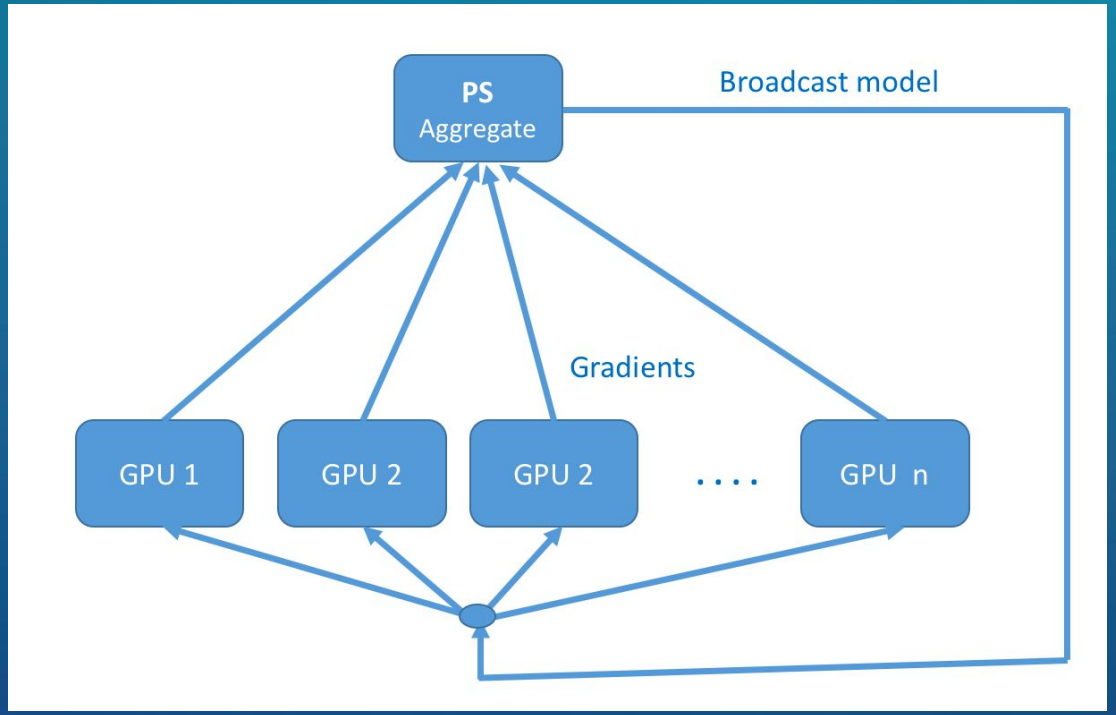

*Image from Towards Data Science

Summary

Productivity matters : teams with better tools/scaling can try out more ideas

SGD
BS = 4096
~1230 sec

→ + LR Scaling
BS *= 16
~140 sec

→ + WARMUP
BS *= 192
~100 sec

→ +LARS
BS *= 192
~110 sec

→ +AMP
BS *= 192
~ 55 sec

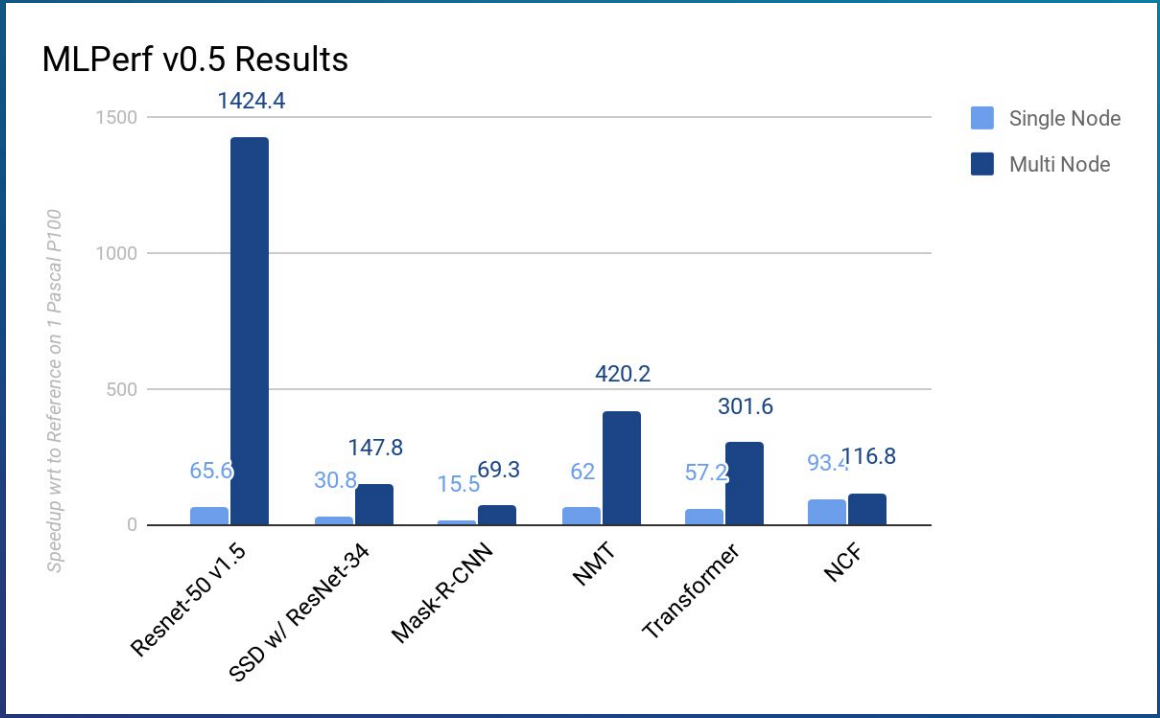

Thank You!

Please remember to complete the session survey in the mobile app.

Follow us @

https://www.linkedin.com/in/alexlinggi/

https://www.linkedin.com/in/swethmandava/

@swethmandava
@mrmendelson

GRACE HOPPER CELEBRATION

ANITA B.ORG

#GHC19